

Chapter 11. Text Mining with R

Reference: Silge, J. and Robinson, D. (2017). Text Mining With R. O'Reilly. Available online at <https://www.tidytextmining.com/>

1. Tidy text format: a table with one token per row.

Token (or Term): a word, a phrase, or several connected words.

Other text data structures:

String: text data is often imported into R as strings (i.e. character vectors).

Corpus: a collection of raw strings annotated with additional meta data and details.

Document-term matrix: a sparse matrix representing a collection (i.e. a corpus) of documents, in which each row stands for a document, each column stands for a term/token, and each entry is, e.g. TFIDF.

`unnest_tokens`: a function in R-package `tidytext` which transform text strings to tidy text format via `data_frame`

```
text1 = c("Long ago, big data was a thick screen, I was here, mainframe computing was  
there", "And now, big data is a thin smart-phone, I am here, cloud computing  
is there", "In future, big data will be tiny particles, I will be here,  
quantum computing will be there")  
# Poem "Big Data" by Professor Yazhen Wang  
> class(text1)  
[1] "character"  
length(text1)  
[1] 3  
> text1  
[1] "Long ago, big data was a thick screen, I was here, mainframe computing was there"  
[2] "And now, big data is a thin smart-phone, I am here, cloud computing is there"  
[3] "In future, big data will be tiny particles, I will be here, quantum computing  
will be there"
```

`text1` is a typical text vector to be analysed. In order to turn it into a tidy text dataset, we need to put it into a data frame using `data_frame` (**NOT `data.frame`!!!**).

```
> library(dplyr)  
> text1_df = data_frame(text1)  
> text1_df
```

```
# A tibble: 3 x 1
  text1
1 Long ago, big data was a thick screen, I am here, mainframe computing was there
2 And now, big data is a thin smart-phone, I am here, cloud computing is there
3 In future, big data will be tiny particles, I will be here, quantum computing ...
```

A tibble is a modern version of data frame. `read_csv` imports data into the tibble format.

```
> install.packages("tidytext")
> library(tidytext)
> unnest_tokens(text1_df, word1, text1)
# A tibble: 48 x 1
  word1
  <chr>
1 long
2 ago
3 big
4 data
5 was
6 a
7 thick
8 screen
9 i
10 was
# ... with 38 more rows
```

```

> text1_tidy=unnest_tokens(text1_df, word1, text1)
> text1_tidy$word1
 [1] "long"      "ago"      "big"      "data"     "was"      "a"        "thick"
 [8] "screen"    "i"        "was"      "here"     "mainframe" "computing" "was"
[15] "there"     "and"      "now"      "big"      "data"     "is"       "a"
[22] "thin"      "smart"    "phone"    "i"        "am"       "here"     "cloud"
[29] "computing" "is"       "there"    "in"       "future"   "big"      "data"
[36] "will"      "be"       "tiny"     "particles" "i"        "will"     "be"
[43] "here"      "quantum"  "computing" "will"     "be"       "there"

```

Output vector word1 discards punctuation, converts the tokens (i.e. words) to lowercase, change smart-phone to smart and phone.

Using pipe: %>%

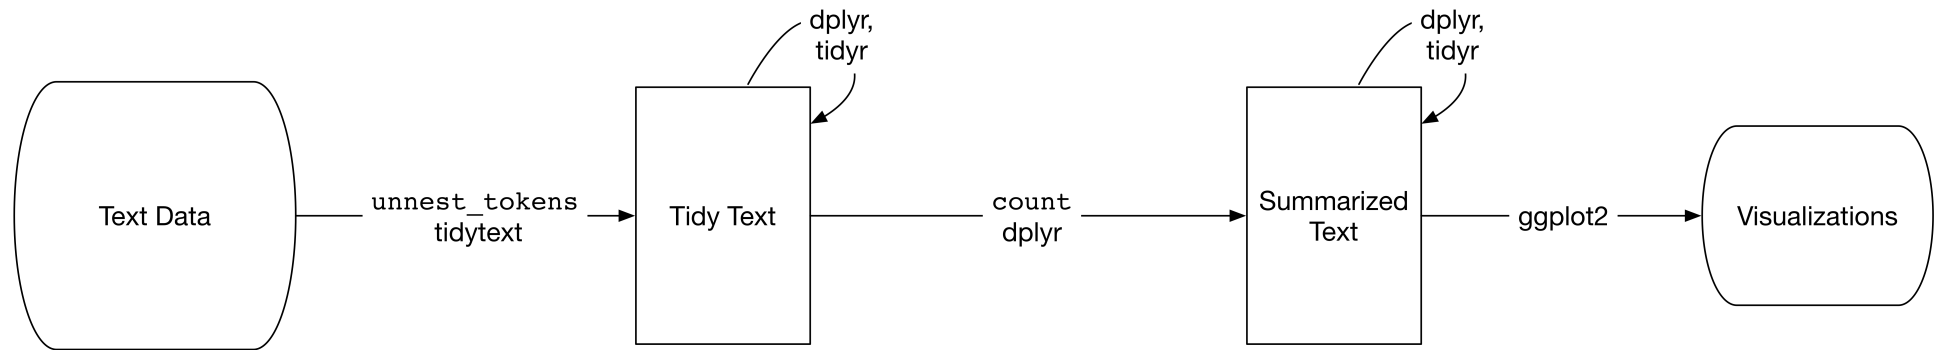
Command `text1_tidy=unnest_tokens(text1_df, word1, text1)` can be equivalently written as

```

> text1_tidy = text1_df %>% unnest_tokens(word1, text1)

```

Data in tidy-text format allow further analysis as illustrated below



For example

```
> count(text1_tidy, word1, sort=T)
# A tibble: 30 x 2
  word1      n
  <chr>    <int>
1 be         3
2 big        3
3 data       3
4 here       3
5 i          3
6 there      3
7 will       3
8 a          2
```

```
9 was          2
10 computing    2
# ... with 20 more rows
```

The above steps can be combined together using pipes:

```
> text1_df %>% unnest_tokens(word1, text1) %>% count(word1, sort=T)
```

Now let us look at the novels by Jane Austen.

```
> install.packages("janeaustenr")
> library(janeaustenr); library(dplyr); library(tidytext)
> prideprejudice[1:11]
[1] "PRIDE AND PREJUDICE"
[2] ""
[3] "By Jane Austen"
[4] ""
[5] ""
[6] ""
[7] "Chapter 1"
[8] ""
[9] ""
[10] "It is a truth universally acknowledged, that a single man
      in possession"
[11] "of a good fortune, must be in want of a wife."
> PP_df <- data_frame(prideprejudice)
> PP_tidy <- PP_df %>% unnest_tokens(word, prideprejudice)
```

Now all the words in *Pride & Prejudice* are in tidy text file PP_tidy.

To load the database of stop words, `data(stop_words)`. Note vector `stop_words` contains all the stop words from 3 lexicons SMART, snowball, onix. To use only the stop words from one lexicon, `stopwords1 = filter(stop_words, lexicon=="SMART")`.

To separate stop words from the others in PP_tidy:

```
> PP_noS <- anti_join(PP_tidy, stop_words)
      # extract non-stop words from PP_tidy
> PP_stop <- semi_join(PP_tidy, stop_words)
      # extract stop words from PP_tidy
```

Now we produce a word-frequency bar-chart using `ggplot2`. It also illustrates the usefulness of piping `%>%`.


```

> library(ggplot2)
> PP_noS %>% count(word, sort=T)
# A tibble: 6,009 x 2
  word      n
  <chr>   <int>
1 elizabeth 597
2 darcy    373
3 bennet   294
4 miss     283
5 jane     264
6 bingley  257
7 time     203
8 lady     183
9 sister   180
10 wickham  162
# ... with 5,999 more rows
> PP_noS %>% count(word, sort=T) %>% filter(n>150)
# A tibble: 13 x 2
  word      n
  <chr>   <int>
1 elizabeth 597
2 darcy    373
3 bennet   294
4 miss     283
5 jane     264

```

6	bingley	257
7	time	203
8	lady	183
9	sister	180
10	wickham	162
11	dear	158
12	collins	156
13	family	151

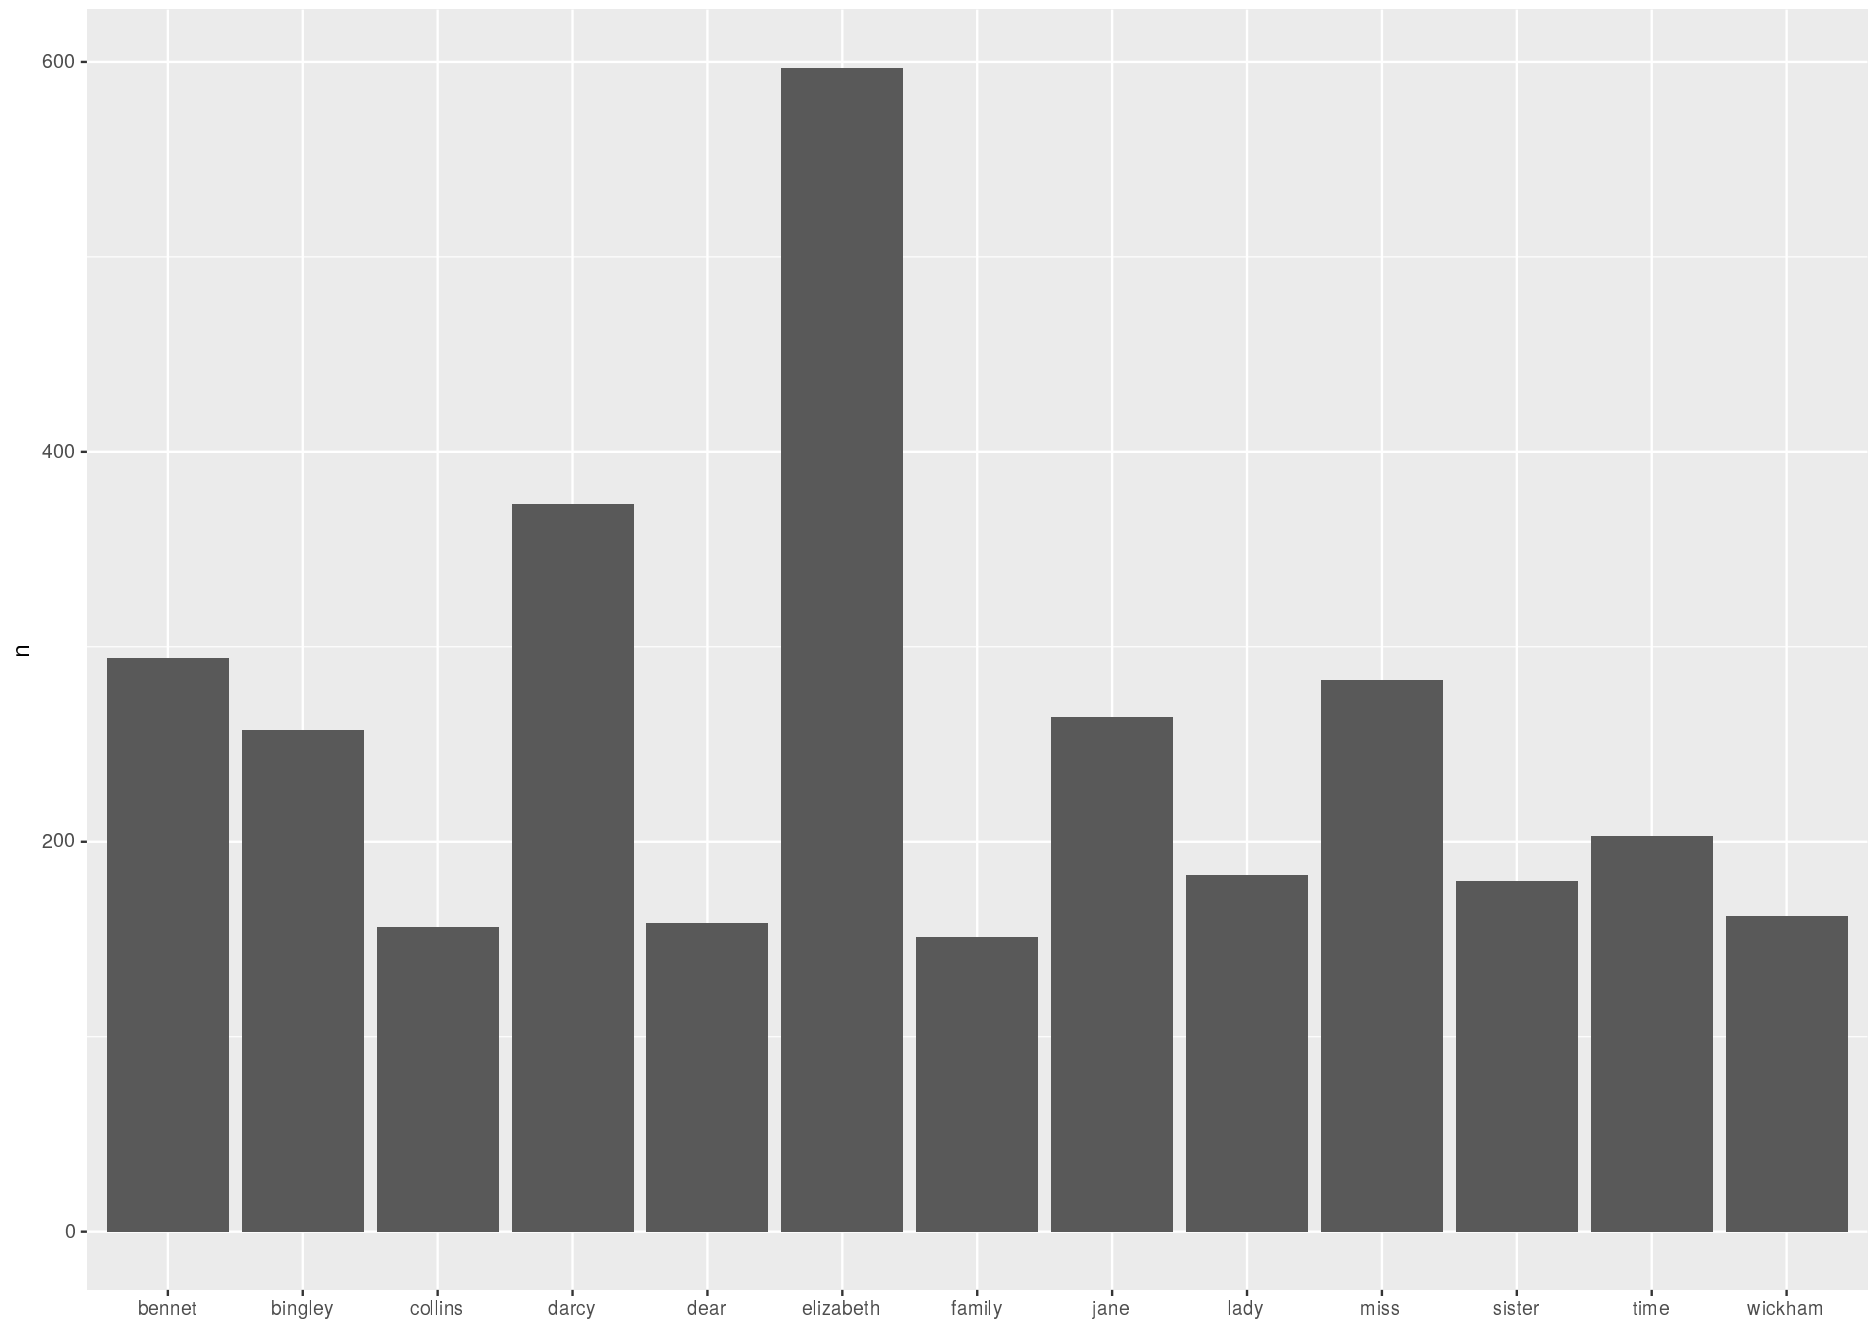
```
> PP_noS %>% count(word, sort=T) %>% filter(n>150) %>% ggplot(aes(word,n)) +  
+ geom_col() +  
+ xlab(NULL)
```

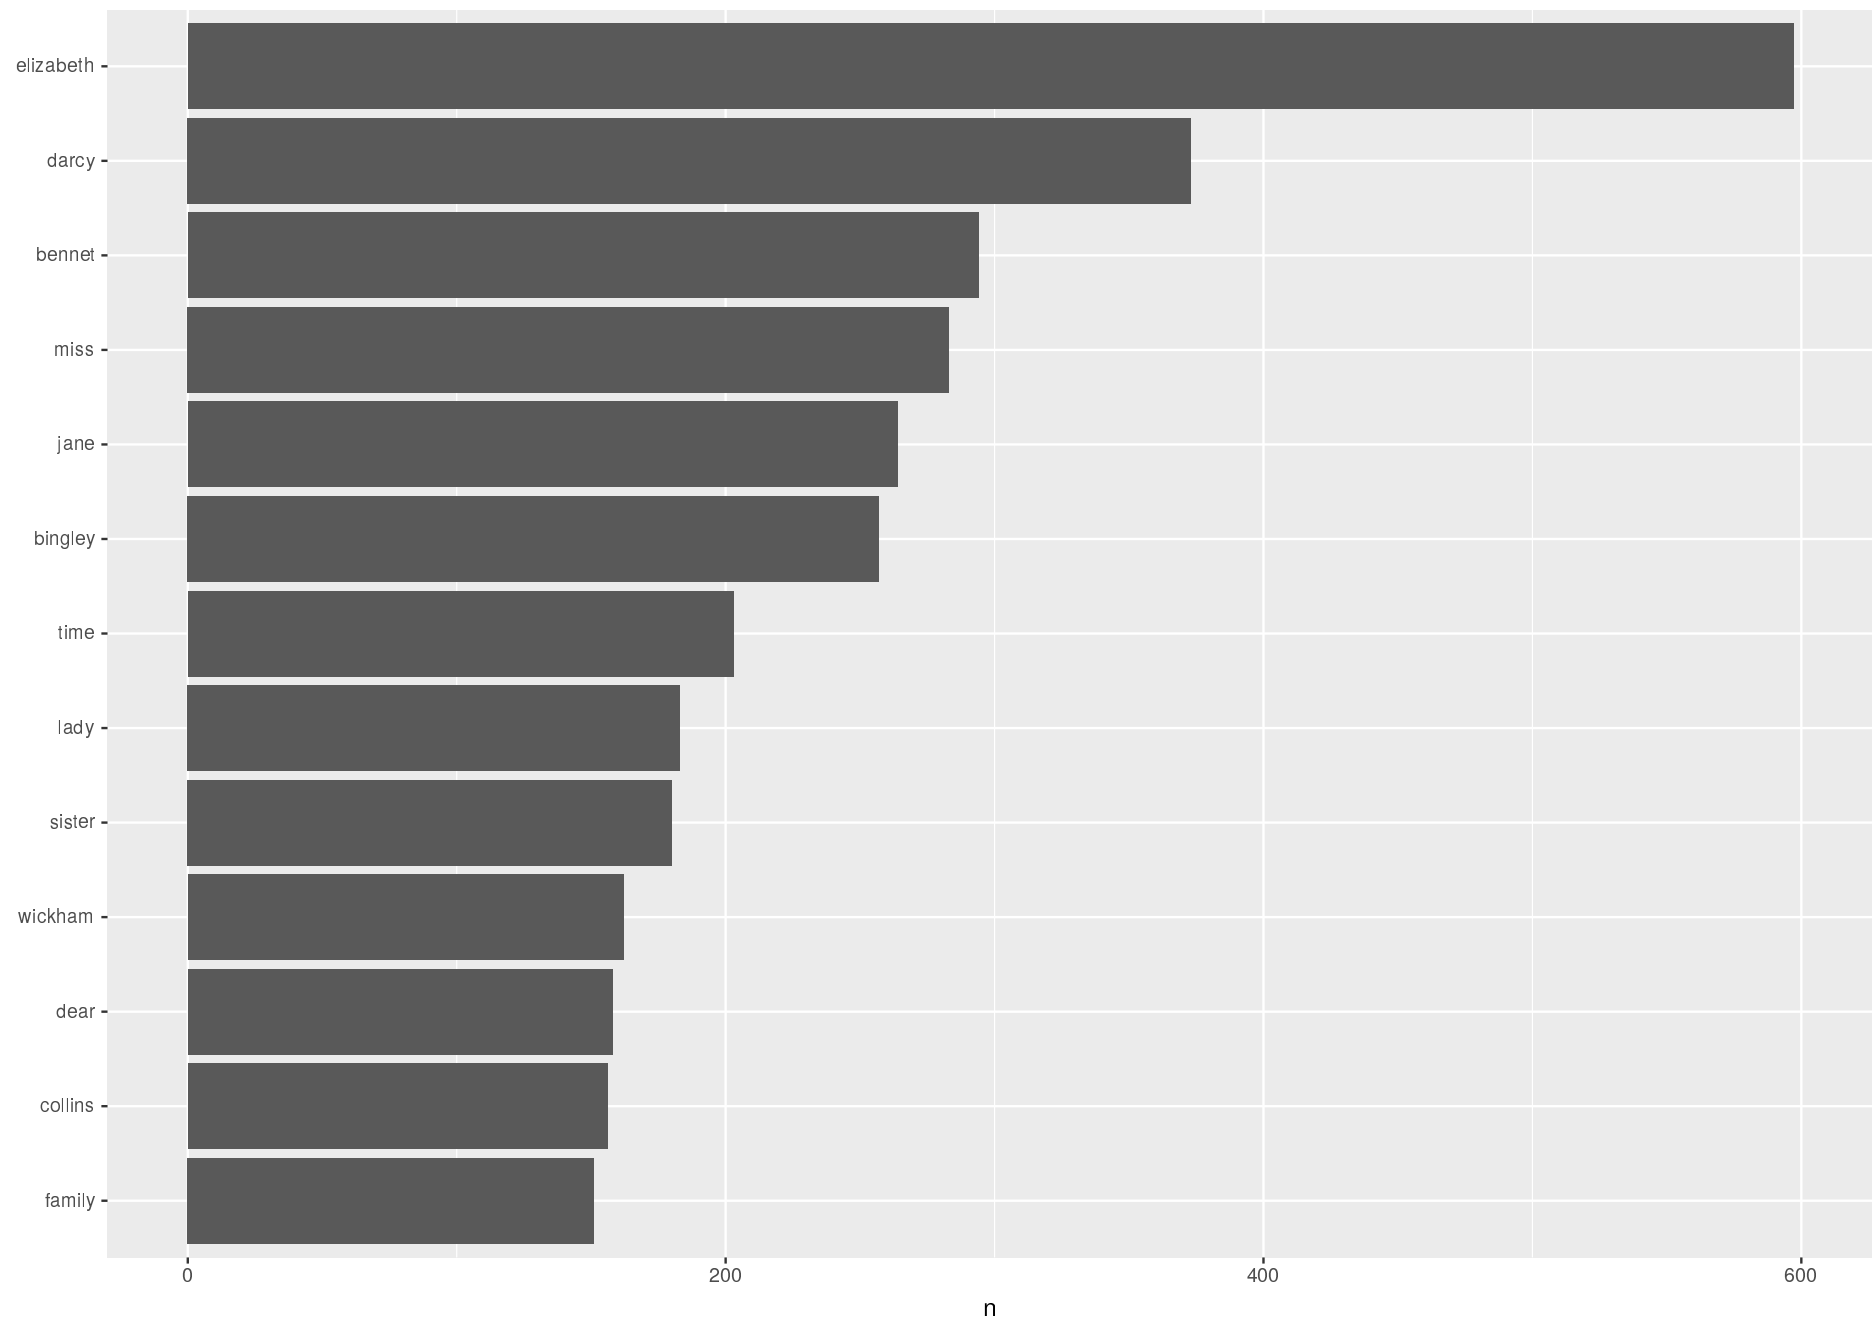
Produce 1st figure

```
> PP_noS %>% count(word, sort=T) %>% filter(n>150) %>% mutate(word=reorder(word,n)) %>%  
+ ggplot(aes(word,n)) +  
+ geom_col() +  
+ xlab(NULL) +  
+ coord_flip()
```

Produce 2nd figure

```
>
```



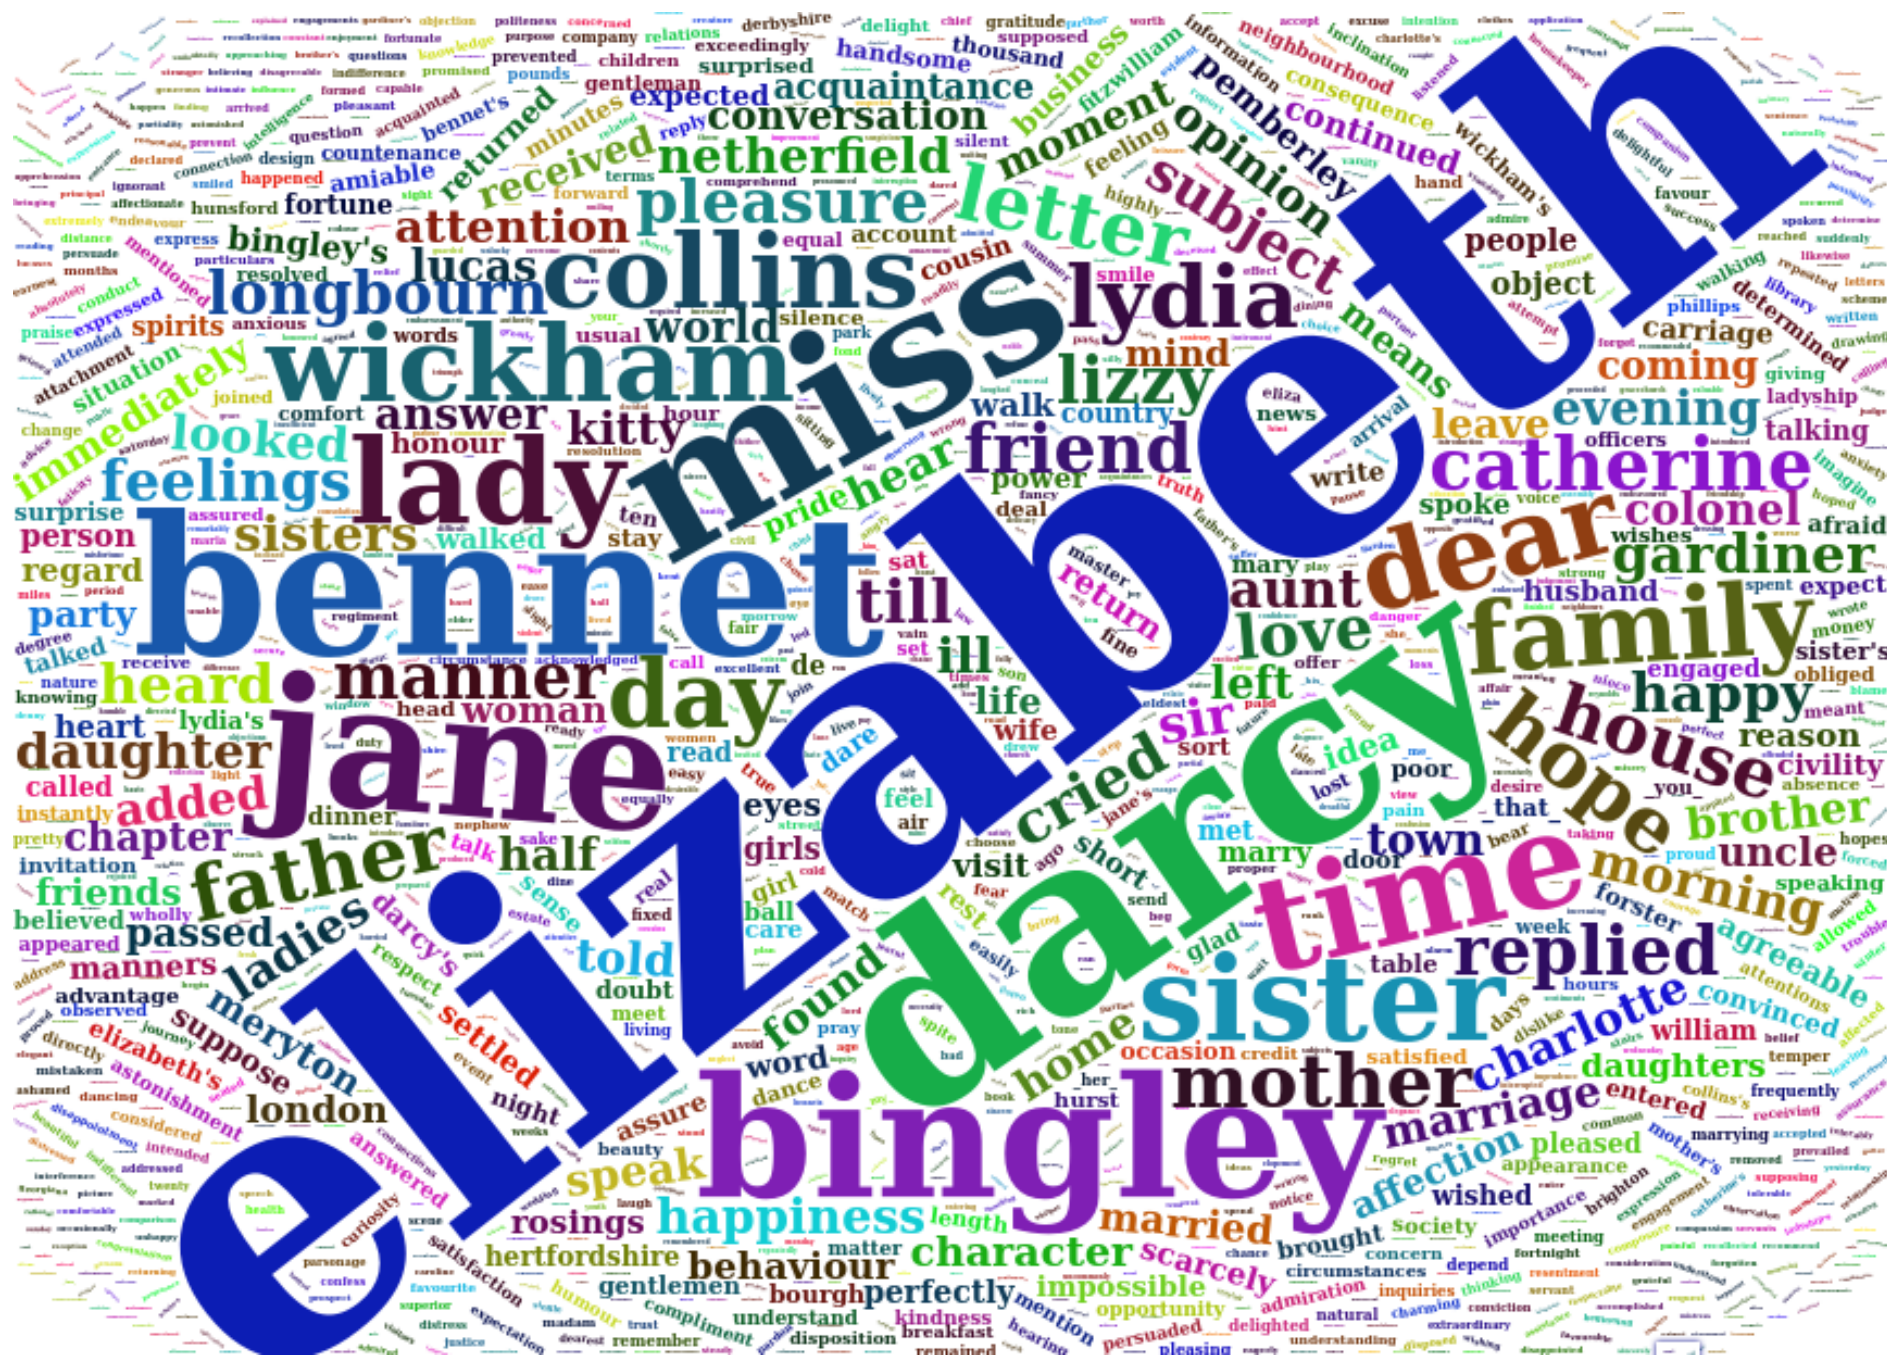


To produce a word cloud plot:

```
> install.packages("wordcloud2")  
> library(wordcloud2)  
> PP_noS %>% count(word, sort=T) %>% wordcloud2()
```

You may also try

```
> PP_noS %>% count(word, sort=T) %>% filter(n>60) %>% wordcloud2()
```



Suppose we try to identify the authorship of a novel. One effective approach is to compare the relative frequencies of stop words in novels.

```
> emma_df <- data_frame(emma)
> emma_tidy <- emma_df %>% unnest_tokens(word, emma)
> emma_stop <- emma_tidy %>% semi_join(stop_words)
> emma_noS <- emma_tidy %>% anti_join(stop_words)
> dim(emma_noS); dim(emma_stop); dim(PP_noS); dim(PP_stop)
[1] 46775      1
[1] 114221     1
[1] 37246      1
[1] 84958      1
```

Surprisingly both novels contain far more stop words than non-stop words.

```
> bind_rows(mutate(PP_stop, book="Pride & Prejudice"), mutate(emma_stop, book="Emma"))
# mutate adds a new column to data.frame
# bind_rows binds data.frames with the same number columns together
# A tibble: 199,179 x 2
  word  book
  <chr> <chr>
1 and   Pride & Prejudice
2 by    Pride & Prejudice
```

```

3 it      Pride & Prejudice
4 is      Pride & Prejudice
5 a       Pride & Prejudice
6 that    Pride & Prejudice
7 a       Pride & Prejudice
8 man     Pride & Prejudice
9 in      Pride & Prejudice
10 of     Pride & Prejudice
# ... with 199,169 more rows

```

```

> bind_rows(mutate(PP_stop,book="Pride & Prejudice"),mutate(emma_stop,book="Emma")) %>%
+ count(book, word)

```

```

# A tibble: 1,056 x 3

```

	book	word	n
	<chr>	<chr>	<int>
1	Emma	a	3129
2	Emma	able	72
3	Emma	about	249
4	Emma	above	12
5	Emma	according	5
6	Emma	accordingly	4
7	Emma	across	7
8	Emma	actually	29
9	Emma	after	161
10	Emma	afterwards	41


```
# ... with 1,046 more rows
```

```
> bind_rows(mutate(PP_stop,book="Pride & Prejudice"),mutate(emma_stop,book="Emma")) %>%  
+ count(book, word) %>% mutate(proportion=n/sum(n))
```

```
# A tibble: 1,056 x 4
```

	book	word	n	proportion
	<chr>	<chr>	<int>	<dbl>
1	Emma	a	3129	0.0157
2	Emma	able	72	0.000361
3	Emma	about	249	0.00125
4	Emma	above	12	0.0000602
5	Emma	according	5	0.0000251
6	Emma	accordingly	4	0.0000201
7	Emma	across	7	0.0000351
8	Emma	actually	29	0.000146
9	Emma	after	161	0.000808
10	Emma	afterwards	41	0.000206

```
> bind_rows(mutate(PP_stop,book="Pride & Prejudice"),mutate(emma_stop,book="Emma")) %>%  
+ count(book, word) %>% mutate(proportion=n/sum(n)) %>% select(-n)
```

```
# A tibble: 1,056 x 3
```

	book	word	proportion
	<chr>	<chr>	<dbl>
1	Emma	a	0.0157
2	Emma	able	0.000361

```

3 Emma about 0.00125
4 Emma above 0.0000602
5 Emma according 0.0000251
6 Emma accordingly 0.0000201
7 Emma across 0.0000351
8 Emma actually 0.000146
9 Emma after 0.000808
10 Emma afterwards 0.000206
# ... with 1,046 more rows

```

Now we use `spread (& gather)` in `tidyr` to put the data in the shape for comparison:

```

> library(tidyr)
> bind_rows(mutate(PP_stop,book="Pride & Prejudice"),mutate(emma_stop,book="Emma")) %>%
+ count(book, word) %>% mutate(proportion=n/sum(n)) %>% select(-n) %>%
+ spread(book, proportion)
# A tibble: 564 x 3
   word      Emma      'Pride & Prejudice'
  <chr>    <dbl>          <dbl>
1 a      0.0157      0.00981
2 able   0.000361    0.000271
3 about  0.00125     0.000613
4 above  0.0000602    0.000105
5 according 0.0000251    0.0000402

```

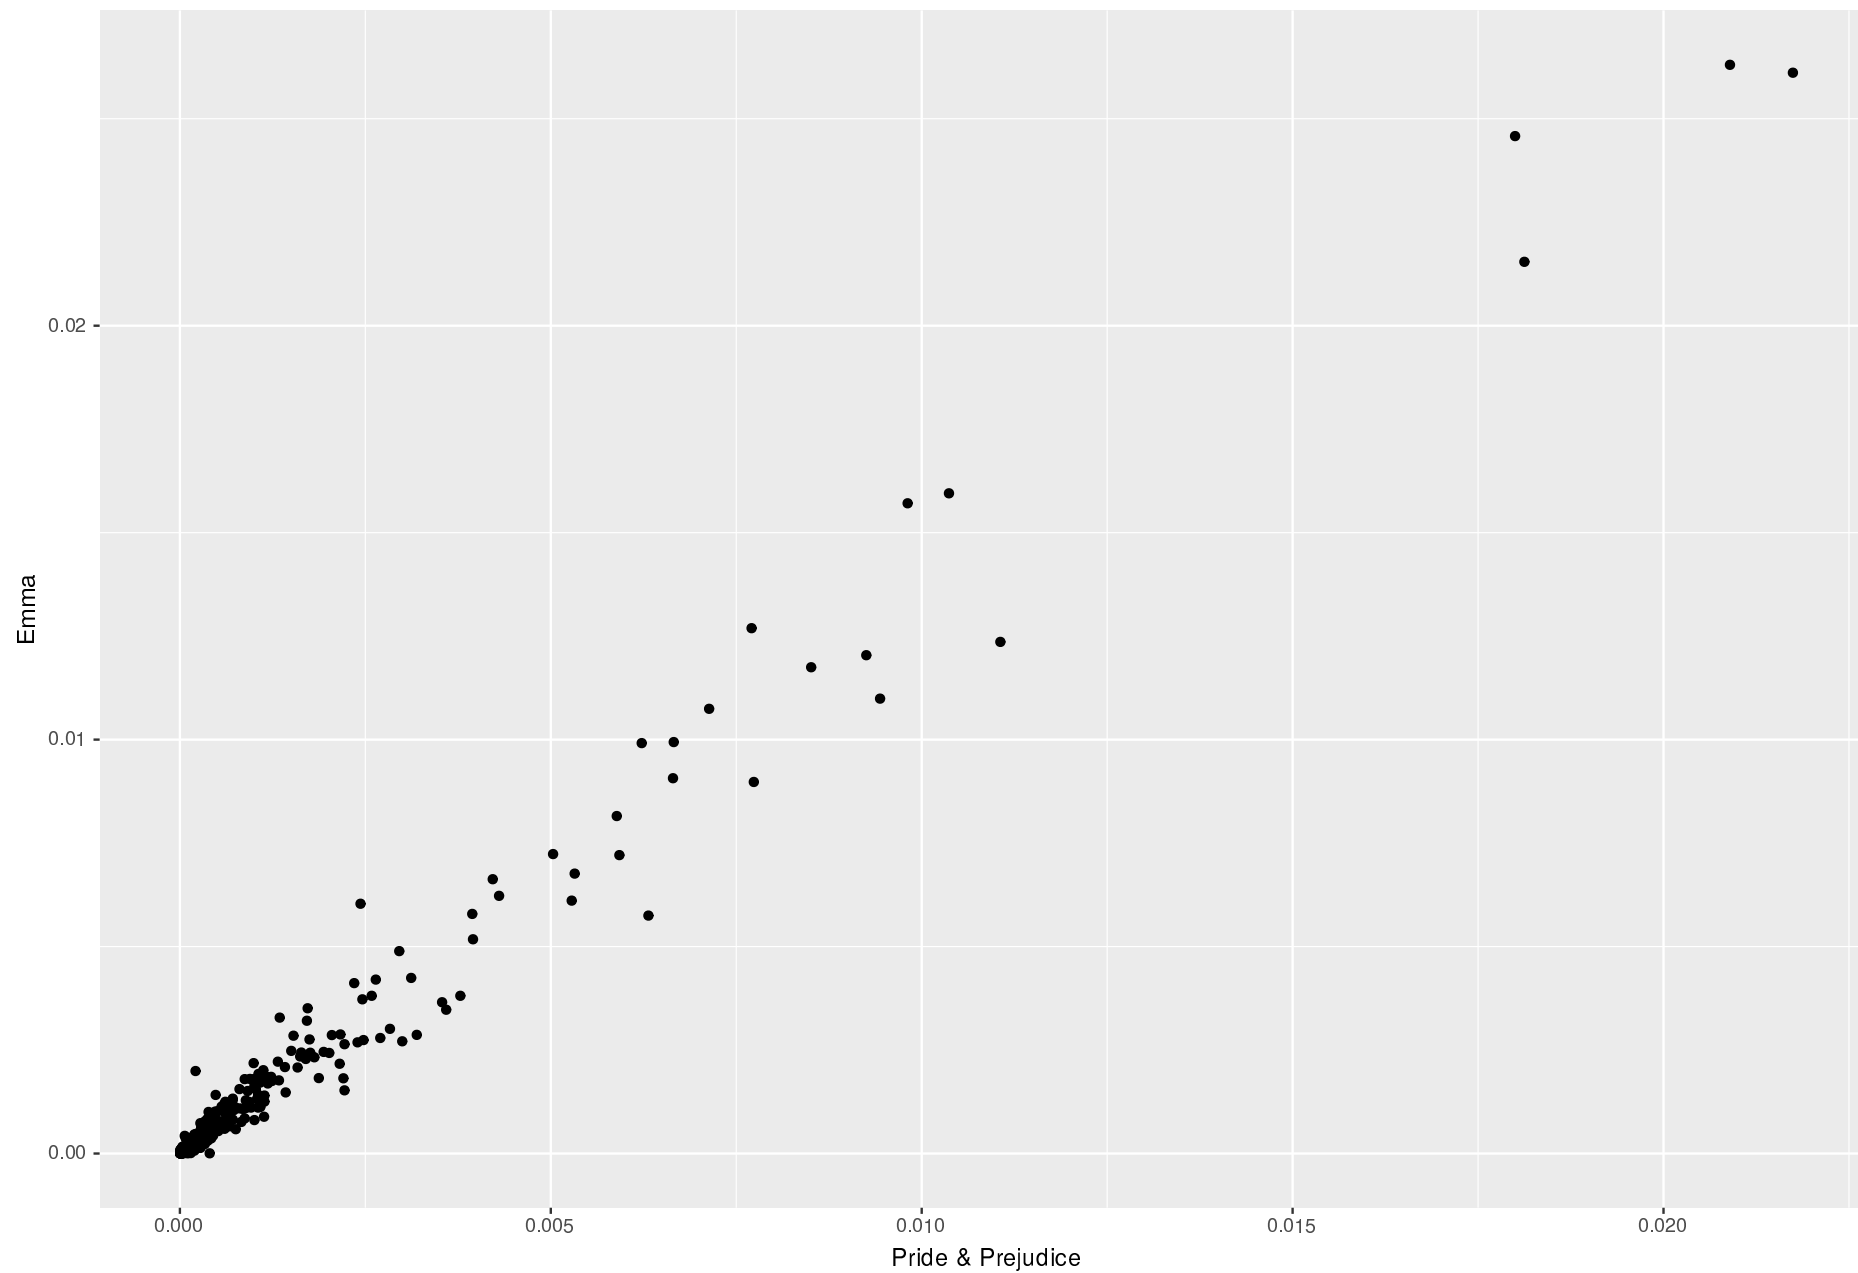
```

6 accordingly 0.0000201          0.0000301
7 across      0.0000351          0.0000251
8 actually    0.000146           0.0000602
9 after       0.000808           0.00100
10 afterwards 0.000206           0.000161
# ... with 554 more rows
> rF = bind_rows(mutate(PP_stop, book="Pride & Prejudice"),
                 mutate(emma_stop, book="Emma")) %>%
+ count(book, word) %>% mutate(proportion=n/sum(n)) %>% select(-n) %>%
+ spread(book, proportion)
> qplot(rF[,3], rF[,2], ylab="Emma", xlab="Pride & Prejudice",
        main="Relative frequencies of stop words in two novels")

```

The figure shows that the relative frequencies of the occurrence of stop words in the two Austen's novels are similar.

Relative frequencies of stop words in two novels



To compare Austen's writings with others, we download 2 Dickens' books from <http://www.gutenberg.org/ebooks/>.

First, Dickens' Great Expectation in html format.

```
> install.packages("rvest") # Package for easy scrape of web pages
> library(rvest)
> GE <- read_html("http://www.gutenberg.org/files/1400/1400-h/1400-h.htm")
> GE_text=GE %>% html_nodes("p") %>% html_text() # Extract text from html file
> GE_df = data_frame(GE_text)
> GE_tidy = GE_df %>% unnest_tokens(word, GE_text)
> GE_stop = GE_tidy %>% semi_join(stop_words)
```

To get Dickens' David Copperfield,

```
> DC = read_html("http://www.gutenberg.org/files/9744/9744-index.htm")
> DC_text = DC %>% html_nodes("p") %>% html_text()
> DC_df = data_frame(DC_text)
> DC_tidy = DC_df %>% unnest_tokens(word, DC_text)
> DC_stop = DC_tidy %>% semi_join(stop_words)
```

Now we combine the relative frequencies of stop words in 4 books together to produce a plot for comparison.

```

> rF4 = bind_rows(mutate(PP_stop, book="Pride & Prejudice"),
  mutate(emma_stop, book="Emma"),
  mutate(GE_stop, book="Great Expectation"),
  mutate(DC_stop, book="David Copperfield")) %>%
+ count(book, word) %>% mutate(proportion=n/sum(n)) %>% select(-n) %>%
+ spread(book, proportion)
> rF4
# A tibble: 659 x 5
  word          'David Copperfield'   Emma   'Great Expectation'   'Pride & Prejudice'
  <chr>          <dbl>          <dbl>          <dbl>          <dbl>
1 a              0.0138          0.00540        0.00698        0.00337
2 able           0.0000707        0.000124       0.0000552      0.0000931
3 about          0.00114          0.000429       0.000552      0.000210
4 above          0.0000966        0.0000207      0.0000552      0.0000362
5 according      0.0000310        0.00000862     0.0000310      0.0000138
6 accordingly    0.0000362        0.00000690     0.00000345     0.0000103
7 across         0.0000948        0.0000121      0.0000759      0.00000862
8 actually       0.0000276        0.0000500      0.0000172      0.0000207
9 after          0.000769         0.000278       0.000504       0.000345
10 afterwards    0.000203         0.0000707      0.0000724      0.0000552
# ... with 649 more rows

> rF4c = rF4 %>% drop_na() # Drop the rows with "na"
> library("GGally", lib.loc=~ /R/x86_64-pc-linux-gnu-library/3.2")
> ggpairs(rF4c[,2:5])

```

